


# Using convolutional neural networks for power consumption prediction

## Exploratory data analysis

The dataset we are working with contains the variable VAR, as well as weather details. Let's have a look at each variable individually and see if we can spot any patterns.



### temp

Let's have a look at temperature (temp), first. Apparently, temperature tends to be the lowest in June, July and August, and the highest in January and February. This means the data was probably collected from a country in the south hemisphere, possibly Australia.

### VAR

This makes sense when we look at power consumption (VAR), which is the highest between June and August – cold months during which heating is required – as well as January and February, likely due to high air conditioning usage. On the other hand, in spring and fall, when the temperatures are mild, power consumption is moderate.

### hum, dewpt, thunder


High humidity (hum) and dew point (dewpt) values as well as thunder frequency are consistent with high temperatures in warmer months, which may suggest humid subtropical climate.

### pressure


There appears to be a data issue with pressure in March. A good way to deal with missing or inaccurate values is to replace them with estimates. This will require inspecting pressure data in more detail but one easy way to deal with the issue would be to use the average pressure value at the same time of the day, from a few days before and after the corrupted value occurred.

For now, let's have a closer look at power consumption (VAR), since it's the value we are modeling.

Zooming into individual months reveals a daily pattern: most likely high usage during daytime and low usage at night. Interestingly, colder months are characterized by a double daily hump. The reason is not clear at this point; one explanation could be that in colder months people consume more power in the morning, as they get ready for work (first daily peak) and in the evening, as they turn the heat up upon returning home (second daily peak). On the other hand, warmer months miss the second peak because the temperature gets lower in the evening anyway, limiting the need for air conditioning.




Looking even closer, at weekly data, here's a good question to ask: Is there a weekly pattern? For example, are weekends different than week days? There is no clear answer yet. There might be a subtle pattern but it's not easy to recognize with the naked eye.




Finally, the daily view reveals how power consumption fluctuates depending on time of day. We notice that the beginning of the day (left side of the chart) starts off quite high. This is leftover usage from the day before, since many adults remain active at least until midnight while entertaining, doing chores, or

simply hanging out with lights switched on. Between 2-6 am the power usage dips for the night, then picks up again in the morning and reaches it's maximum around noon (middle of the chart). From there, it stays at a relatively high level until midnight (right side), as household members go about their day. This hourly pattern is very strong and will play a significant role in determining VAR values throughout the day.



A quick glance at the correlation matrix sheds more light on our knowledge/findings so far. First of all, there is a strong negative correlation between power consumption and working/non working days (Day and Holiday variables). Specifically, power consumption is lower during weekends and bank holidays. On the other hand there is no direct correlation between power consumption and temperature, despite our earlier observations. This is because the relationship between the two variables is indirect; it's strongly affected by the daily and weekly pattern.



Problems like this, with intricate variable relationships, where if-then rules are hard to grasp intuitively, are good candidates for deep learning models.

## Convolutional Neural Networks

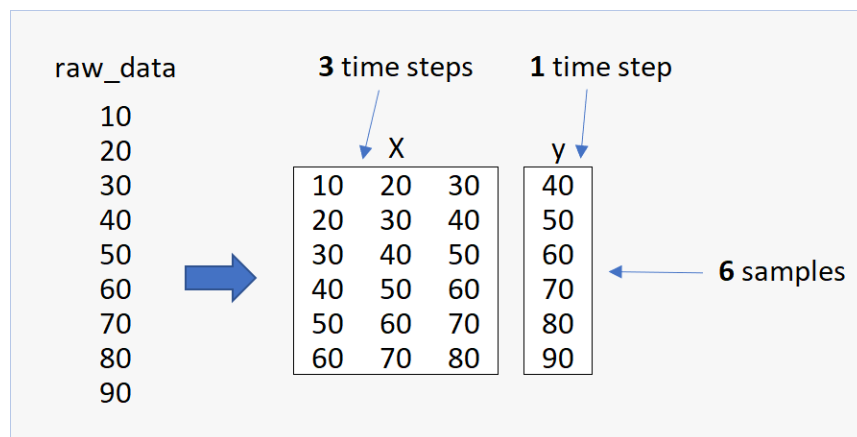
Traditionally, LSTM (Long short-term memory) networks are used for modeling time series. In this exercise, however, we will explore CNN, or Convolutional Neural Networks.

### Data transformation

Before we get started, we need to transform our dataset in a way that a CNN can understand it. Specifically, CNN calls for a three-dimensional dataset, which consists of so called samples, time steps, and features. To develop an intuition on how this three-dimensional structure might look, let's first translate the time series problem into a supervised machine learning problem.

Imagine a simple string of numbers, 10, 20, 30 and so on, until 90. How can we represent it as a classical supervised machine learning problem with an **X** matrix containing features, and a **y** vector containing the target variable? One way would be to take each number, or observation, and notice the values that came before it.

For example, we know that before 90, there is 60, 70 and 80. We could choose to model this relationship as 60, 70 and 80 being our input, and 90 – output data. To use CNN terminology, we have just constructed a single **sample**, with 3 input **time steps** and 1 output **time step**. But here's the tricky part. The three time steps in the input, are a single **feature**. Remember we decomposed a single time series into multiple steps, but these steps are still part of a single feature.



Here's Python code that completes the transformation:

```

# split a univariate sequence into samples
def split_sequence(sequence, n_steps):
    X, y = list(), list()
    for i in range(len(sequence)):
        # find the end of this pattern
        end_ix = i + n_steps
        # check if we are beyond the sequence
        if end_ix > len(sequence)-1:
            break
        # gather input and output parts of the pattern
        seq_x, seq_y = sequence[i:end_ix], sequence[end_ix:]
        X.append(seq_x)
        y.append(seq_y)
    return array(X), array(y)

```

```

# define input sequence
raw_seq = [10, 20, 30, 40, 50, 60, 70, 80, 90]
# choose a number of time steps
n_steps = 3
# split into samples
X, y = split_sequence(raw_seq, n_steps)
# summarize the data
print("X shape: " + str(X.shape))
print("y shape: " + str(y.shape))
for i in range(len(X)):
    print(X[i], y[i])

```

```

X shape: (6, 3)
y shape: (6,)
[10 20 30] 40
[20 30 40] 50
[30 40 50] 60
[40 50 60] 70
[50 60 70] 80
[60 70 80] 90

```

Now, we need to reshape our data into a three-dimensional structure.

```

# reshape from [samples, timesteps]
# into [samples, timesteps, features]
n_features = 1
X = X.reshape((X.shape[0], X.shape[1], n_features))
print(X.shape)
print(y.shape)
for i in range(len(X)):
    print(X[i], y[i])

```

```

(6, 3, 1)
(6,)

```

Finally, we are ready to define a CNN model using Keras library.

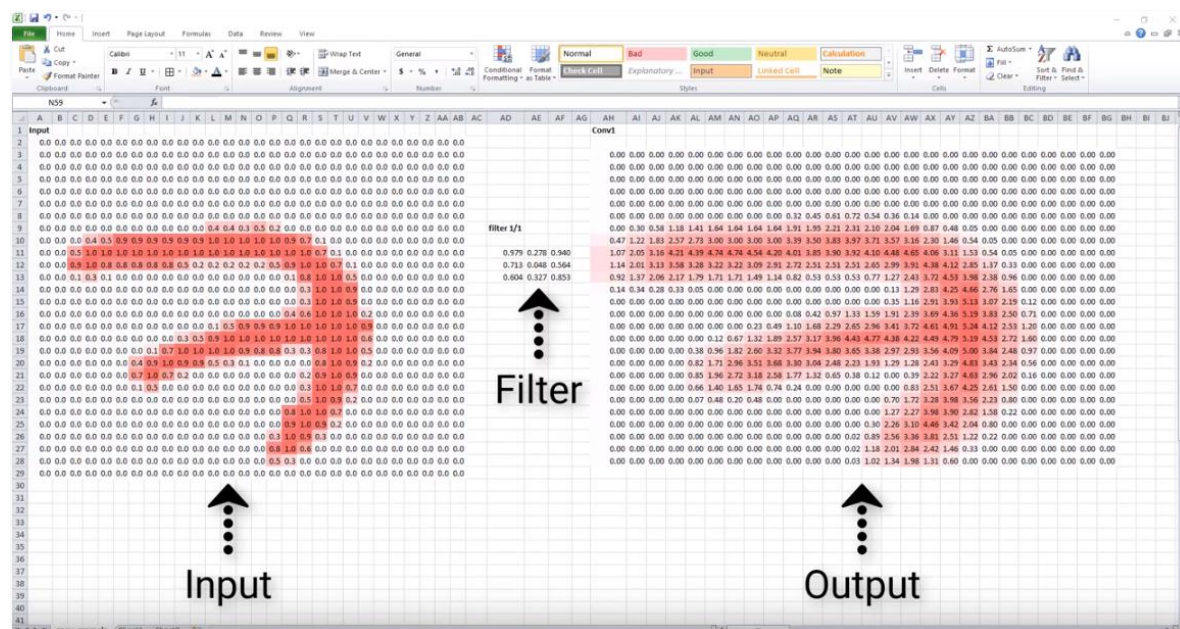
## CNN refresher

If you need a refresher on what a convolutional layer does, Jeremy Howard, the author of fast.ai course, has created a great spreadsheet that implements the logic.

In short, CNN's are traditionally used for image recognition. We represent an image as a matrix of numbers, where each number denotes pixel intensity. The darker the color, the higher the value (let's stick with grayscale for now).

We are going to generate a so called **filter**. You can think of a filter as a small window of random numbers, that slides across the image, computing the dot product between the pixel intensity values and the filter values. We will then save the result as a new image.

Notice that the resulting image is not very different from the original one. That's because the background of the original one is mostly white, with pixel intensity equal to 0. Multiplying zeros by random numbers still gives us zeros; it is around the edges of the shape, where the magic happens. The filters can pick up on various characteristics of the shape and may soften or sharpen the curves, or blur or enhance the edges and so on.

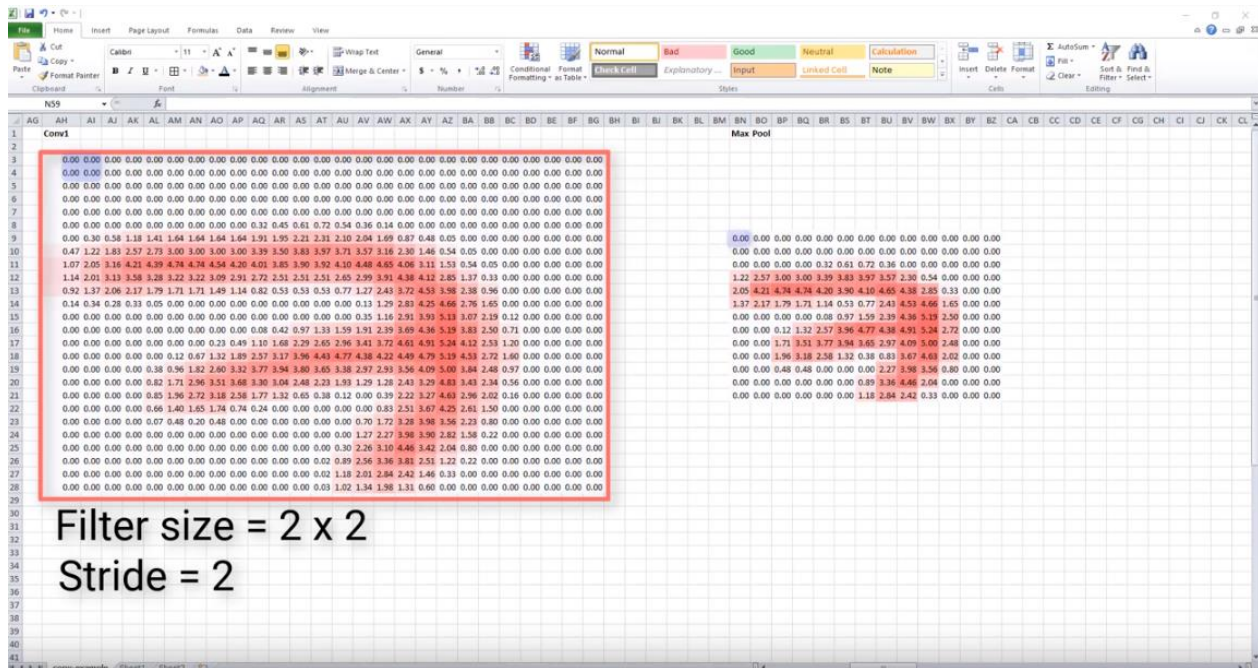


Source: <https://www.youtube.com/watch?v=V2h3IOBDvrA>

In our model, we will generate 64 filters, of the size 2x2. (Jeremy's filter is 3x3)

The second layer we create is a **pooling** layer. Think of pooling as a compression mechanism that makes the image smaller, while preserving its key characteristics. Technically, pooling is achieved by sliding a filter across the image once again but this time, instead of multiplication, we are simply saving the maximum pixel value from the filter's region.





Next, we flatten the data, and create a fully connected, dense layer. We are using **relu** activation function, which replaces negative values with zeros. The efficient **adam** optimizer adjusts the learning rate so we don't have to define it in trial and error, and finally, we used the **mean squared error** as our loss function. We fit the model for 1000 epochs, then ask it to predict the value associated with 70, 80 and 90. The answer we're given is 100.5 – not bad!

```
# define model
model = Sequential()
model.add(Conv1D(filters=64, kernel_size=2, activation='relu',
                 input_shape=(n_steps, n_features)))
model.add(MaxPooling1D(pool_size=2))
model.add(Flatten())
model.add(Dense(50, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mse')
# fit model
model.fit(X, y, epochs=1000, verbose=0)
# demonstrate prediction
x_input = array([70, 80, 90])
x_input = x_input.reshape((1, n_steps, n_features))
yhat = model.predict(x_input, verbose=0)
print(yhat)

[[100.502815]]
```

This concludes our very simple, naïve example.

Let's move on to some more interesting scenarios - let's imagine three time series as follows.

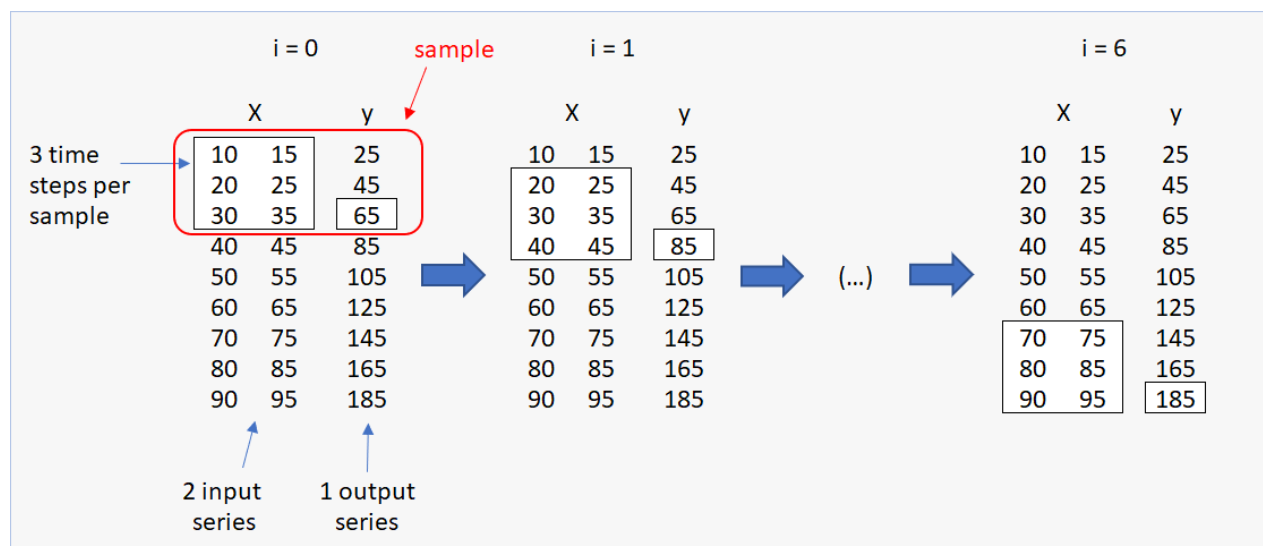
X		y
10	15	25
20	25	45
30	35	65
40	45	85
50	55	105
60	65	125
70	75	145
80	85	165
90	95	185

Our goal is to predict the value of  $y$ , given  $X$ . For instance, given that  $X = 95$  and  $105$ , what will be the value of  $y$ ?

A keen eye will observe that  $y$  is the sum of  $X$ .  $10 + 15 = 25$ .  $20 + 25 = 45$  and so on, so  $90 + 95 = 195$ . Of course, in real life, variable relationships can be much harder to spot, and we want to construct a reliable mechanism that will predict that  $95$  and  $105$  are associated with the number  $195$ , without knowing the nature of the relationship.

We could use each row as a single observation and try to predict the outcome, but this approach doesn't take advantage of the temporal structure of the data. So let's expand the concept of an observation from a single row to, say, three.

We will build our first **sample** that associates the number  $65$ , with not only  $30$  and  $35$ , but also the two preceding pairs,  $20$  and  $25$ , as well as  $10$  and  $15$ . The input of our sample will therefore consist of two **time series** and three **time steps**.



To construct more samples, we simply slide the window down, until we reach the end of the dataset. We will then have created 7 samples.



The dimensionality of our CNN-ready input dataset is now [7, 3, 2]. That is 7 samples, 3 time steps, 2 features. We can build a model in the exact same way.

```
# multivariate data preparation
from numpy import array
from numpy import hstack

# split a multivariate sequence into samples
def split_sequences(sequences, n_steps):
    X, y = list(), list()
    for i in range(len(sequences)):
        # find the end of this pattern
        end_ix = i + n_steps
        # check if we are beyond the dataset
        if end_ix > len(sequences):
            break
        # gather input and output parts of the pattern
        seq_x, seq_y = sequences[i:end_ix, :-1], sequences[end_ix-1, -1]
        X.append(seq_x)
        y.append(seq_y)
    return array(X), array(y)

# choose a number of time steps
n_steps = 3
# convert into input/output
X, y = split_sequences(dataset, n_steps)
# summarize the data
print("X shape: " + str(X.shape))
print("y shape: " + str(y.shape))
for i in range(len(X)):
    print(X[i], y[i])
```

```
# define model
model = Sequential()
model.add(Conv1D(filters=64, kernel_size=2, activation='relu',
                 input_shape=(n_steps, n_features)))
model.add(MaxPooling1D(pool_size=2))
model.add(Flatten())
model.add(Dense(50, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mse')
# fit model
model.fit(X, y, epochs=1000, verbose=0)
# demonstrate prediction
x_input = array([[80, 85], [90, 95], [100, 105]])
x_input = x_input.reshape((1, n_steps, n_features))
yhat = model.predict(x_input, verbose=0)
print(yhat)
```

```
[[206.07697]]
```